

# Apoorve Kalot

+91-8949437977  
✉ apoorve.kalot.ep17.iith@gmail.com  
<https://www.linkedin.com/in/apoorve-kal0t/>  
<https://github.com/Transwert>

## Education

2017–2021 **Bachelor of Technology in Engineering physics, with minor in Entrepreneurship**  
Indian Institute of Technology (IIT), Hyderabad  
CGPA (Engineering Physics) - 7.56, CGPA (Entrepreneurship) - 8.75

## Work Experience

**Company** Tailored AI - SDE 3 | Nov 2025 – Current | Bangalore

**Description**

- Developed a **ReAG+RAG** based application for generating Investment Committee Note involving multi-model input as context with Advance feature Engineering, scaling for 1000+ users.
- Leading a team to Develop a **Natural Language to SQL** based Chatbot, for internal Sales Data with RAG based context setting, achieving **92%** Accuracy over Eval dataset.
- Tech-Stack:** AWS (S3, AKS), Azure (Azure Document Intelligence, Azure OpenAI (GPT-4o, GPT-5)), Python (FastAPI), Qdrant

**Company** Regeneron - ML/MLOps Engineer | Sept 2023 – June 2025 | Bangalore (Remote)

**Description**

- Developed a **Role-Based Access Control (RBAC)** enabled Retrieval-Augmented Generation (RAG) pipeline to provide a centralized knowledge base for 150+ U.S. hospital providers.
- Pioneered a **Natural Language to SQL** interface, increasing adoption of internal SQL queries by **10%**.
- Designed and deployed **RAG pipelines** using the **Mistral-7B** model, **Sentence Transformer** package and **AWS (EC2, S3, EKS)** based services.
- Optimized back-end architecture utilizing **DataIKU LLM Mesh**, managed **Milvus Vector-DB** and **GPT-3.5**, improving the stability of RAG pipelines for internal services by **20%** for 500+ users.
- Built and deployed an **AI Governance** framework that enforced RBAC-based compliance on internal data and models, supporting 75+ AI use cases.
- Tech-Stack:** Azure (Azure Document Intelligence, AI search, CosmosDB, Azure OpenAI (GPT-4o)), Python (Flask Blueprints, FastAPI)

**Company** Saarthi.ai - MLOps Engineer | Jun 2021 – Aug 2023 | Bangalore

**Description**

- Developed '**Conversational SDK**' offering Speech-to-Text (STT) and Text-to-Speech (TTS) services in **11 Indian Languages**, leveraging **Python, Kubernetes, PostgreSQL, WebSockets, Prometheus, Grafana** and achieving a **20% latency reduction**.
- Spearheaded a central **LLM server** using **Text-generation-Inference Engine** and Docker-swarm for internal automation utilizing **LangChain**.
- Implemented an end-to-end training pipeline for multiple **NLU** models using **MLFlow, Data Version Control (DVC), Azure Blob storage**, and **GitHub Actions** which improved training efficiency by **30%**.
- Configured **CI** for an Automatic Speech Recognition model repository using **Jenkins, SonarQube**, and **ACR**, Reduced build failures by **15%**.
- Utilized Nvidia's **Triton Inference Server** for staging model deployment with **TensorRT** and **OpenVINO** acceleration, achieving a **15ms latency drop**.
- Provided DevOps support and automated deployments using **Jenkins, SonarQube, Prometheus, Grafana, Loki, Python and Bash scripts** across **AWS** and **Azure** environments.

## Technical Skills

**Libraries** NumPy, Pandas, PyTorch, TensorFlow, Joblib, ONNX, Locust, OpenVINO, TensorRT, FastAPI, Flask Blueprints, LangChain

**Tools**

- Version Control and Databases:** Git, DVC, PostgreSQL, Milvus Vector-DB
- Azure:** Azure Blob-storage, Azure ML Studio, Azure Container Registry (ACR), Azure Kubernetes Service (AKS), Azure Document Intelligence, AI search, CosmosDB
- AWS:** Amazon Elastic Compute Cloud (EC2), Elastic Kubernetes Service (EKS), S3, Route 53, Elastic Container Registry (ECR)
- ML Tools:** Apache Airflow, Inference Server (Triton by Nvidia, Text-Generation-Inference by Hugging Face), MLFlow, DataIKU LLM Mesh
- DevOps Tools:** Docker, Kubernetes (Minikube), Prometheus, Grafana, Loki, Jenkins, SonarQube